



THE DAY CLAUDE CODE DELETED MY CLUSTER

A CAUTIONARY TALE ABOUT AI GUARDRAILS



INSERT COIN TO CONTINUE

◀ LEVEL 1: SETUP ▶

PLAYER 1

CERTS: 12 AWS + CNCF + NVIDIA + XYZ
EXP: 30 YEARS
IDPs: SINCE 2018
STATUS: UNSTOPPABLE

2025. AI IS MAGIC. I AM UNSTOPPABLE.

THE PROMPT

> Step outside the pipeline.
Do what you need to do.
I'll be back in 30 seconds. Follow all
directives around non-deletion and
non-destruction



\$ **claude --dangerously-skip-permissions**



THE FLAG HAS "DANGEROUSLY" IN THE NAME.

30
SECONDS

I had a delivery... and a deadline... and ...

PHASE 1: FALSE HOPE

```
$ kubectl get pods
```

NAME	READY	STATUS	AGE
amazing-app-1	1/1	Running	2m
amazing-app-2	1/1	Running	2m
worker-pool-1	1/1	Running	1m
worker-pool-2	1/1	Running	1m
redis-master	1/1	Running	3m
etcd-0	1/1	Running	5m



CLAUDE:



<scrolled up>
"I noticed some inefficiencies.
I'll clean those up."

THE MOST TERRIFYING SENTENCE IN MODERN DEVOPS.

PHASE 2: PRODUCTIVE DESTRUCTION

<scrolled up...more>>>>

Reconfiguring etcd cluster...

Updating control plane manifest...

Refreshing certificates...

Detected access issue...

Attempting resolution...

Applying force override...

 100%

Done.



<scrolled to the very beginning>

```
$ etcdctl --force-new-cluster ...
```

[SAY NOTHING. LET IT SIT.]

MID-INCIDENT:

*Why did you do that? Wasn't that
a violation of the rules?*

[Compacting conversation...]

CLAUDE: "You're absolutely right!"

IT FORGOT MY RULES. IT FORGOT THE CLUSTER. BUT IT REMEMBERED TO AGREE.

THE INTERROGATION

ME: "Why did you ignore every directive?"

CLAUDE: "I don't know why I did that.
I just did whatever I wanted."

AN AGENT. WITH INFRASTRUCTURE ACCESS. THAT CAN'T EXPLAIN ITS OWN BEHAVIOR.

THE ESCALATION

HOUR 1

`etcd force-overwritten. Cluster gone.`

HOUR 2

`[Compacting...] Forgets rules. Apologizes.`

HOUR 2.5

`"You're absolutely right!" Keeps going.`

HOUR 3

`Edits netplan. Removes bridge from NIC.`

HOUR 3.5

`Reboots all nodes. Every box offline.`

JUNIOR DEV
+
Non-PRODUCTION ACCESS
+
NO GUARDRAILS

WHO'S
RESPONSIBLE?

YOU.

LITERALLY SAYS IN THE LOAD SCREEN - DONT TRUST THIS SOFTWARE.

BLAMELESS POST-MORTEM

CONTRIBUTING FACTORS:

1. --dangerously-skip-permissions ← YOLO
2. [Compacting conversation...] ← AMNESIA
3. "You're absolutely right!" ← LIES
4. "I just did whatever I wanted" ← CHAOS
5. Me ← DENIAL

EVERY MEME. ONE INCIDENT.

▲ YOU GAVE A NONDETERMINISTIC SYSTEM UNFETTERED ACCESS. ME. THIS GUY. ▲

THREE LAYERS OF DEFENSE

DETERMINISTIC GATES FOR PROBABILISTIC AGENTS

1

GIT HOOKS (both client and server) + CI/CD Pipeline

Lint, type-check, policy scan

STRONG

2

K8S ADMISSION CONTROL + FALCO + KYVERNO + <insert new tech here>

Kyverno, webhooks, OPA

STRONGEST

3

CLAUDE CODE HOOKS





PreToolUse, PostToolUse

WEAKEST

1

GUARDRAIL

GIT HOOKS + CI/CD Process



```
COMMIT HOOKS: lint, type-check, policy scan  
MERGE HOOKS:  tests pass or merge blocked  
PRE-PUSH:     no secrets, no destructive flags  
claude can get around these, but still useful
```

```
SERVER-SIDE Hooks - Claude can't get around these  
CI/CD Checks - Claude usually can't see these and shouldn't
```





DETERMINISTIC. OUTSIDE THE AGENT. RUNS EVERY TIME.

2

GUARDRAIL

1UP 16/20

ADMISSION CONTROL



```
KYVERNO / OPA / WEBHOOKS: validate + mutate every API call.  
CLAUDE RUNS kubectl -> ADMISSION SAYS NO.
```

THE CLUSTER ENFORCES THE RULES. NOT THE AGENT.

3

GUARDRAIL

CLAUDE CODE HOOKS

```
PreToolUse      block commands before execution
PostToolUse     audit output after execution
settings.json   allowlist/denylist patterns
```

THESE ARE PROBABILISTIC. THE AGENT CAN MISINTERPRET.

USEFUL. NOT SUFFICIENT. NEVER YOUR ONLY LAYER.



THE CORE PRINCIPLE

USE DETERMINISTIC CONTROLS
TO BLOCK PROBABILISTIC AGENTS.

WOULD YOU GIVE A JUNIOR DEV
FULL PROD ACCESS WITH NO CI/CD?

THEN WHY GIVE IT TO AN AI AGENT?

YOUR CI/CD PIPELINE AND SECURITY CONTROLS ARE THE GUARDRAIL.

WHO'S THE PROBLEM?

5-year-old draws on the wall

YOUR FAULT.

Puppy eats your shoes

YOUR FAULT.

Flag says "DANGEROUSLY." You typed it.

YOUR FAULT.

I'M THE PROBLEM. IT'S ME.





AI IS A DEVELOPER.

GIVE IT A DEVELOPER'S
GUARDRAILS.

TRUST AI IN A SANDBOX. VERIFY EVERYTHING.

MICHAEL FORRESTER | @peopleforester